

Using GPUs to improve the computational efficiency of Gaussian process models

Colin Rundel

Duke University

February 27, 2014

1 Background

2 Migratory Bird Spatial Assignment Model

3 Speciated PM_{2.5} Modeling

4 Related and Future Work

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample?

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood?

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Fitting a model?

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Fitting a model? - Update $\boldsymbol{\Sigma}$

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Fitting a model? - Update $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^2)$

The problem with GPs ...

Unless you are lucky (or clever), Gaussian process models are difficult to scale to large problems.

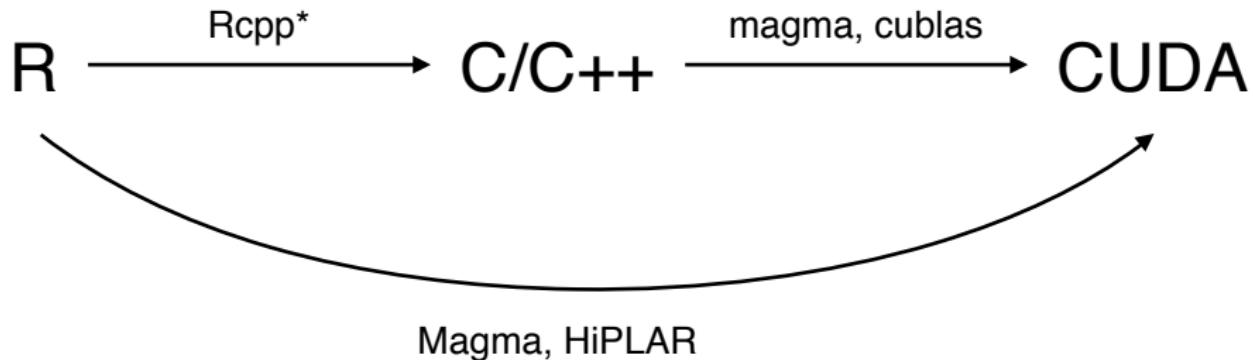
MVN Density:

$$f(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

- Drawing a sample? - Choleksy of $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Evaluating the likelihood? - Invert $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^3)$!
- Fitting a model? - Update $\boldsymbol{\Sigma}$ - $\mathcal{O}(n^2)$

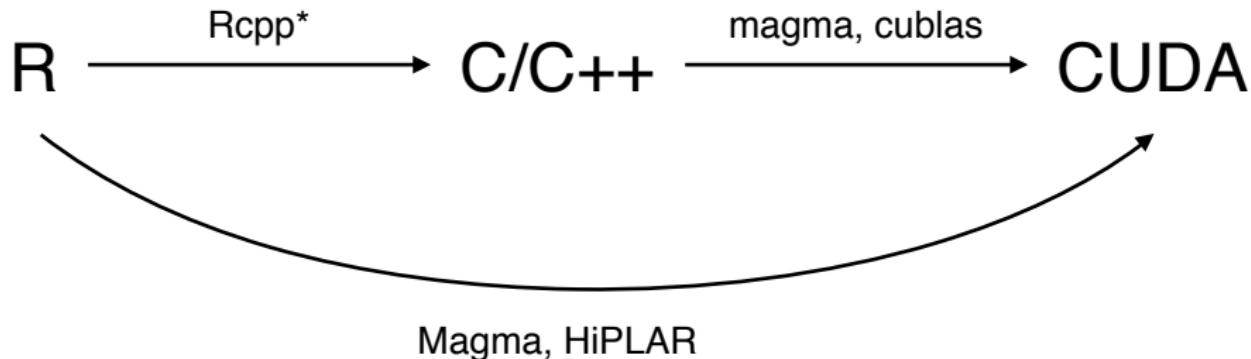
I don't have any clever solutions for these problems, but there are very smart people out there who are working on them.

Tools and Optimization



Regardless of tools or workflow, measuring / profiling performance is critical.

Tools and Optimization



Regardless of tools or workflow, measuring / profiling performance is critical.

"Data or it didn't happen."

1 Background

2 Migratory Bird Spatial Assignment Model

3 Speciated PM_{2.5} Modeling

4 Related and Future Work

Background

Using intrinsic markers (genetic and isotopic signals) for the purpose of inferring migratory connectivity.

- Existing methods are too coarse for most applications
- Large amounts of data are available (>150,000 feather samples from >500 species)
- Genetic assignment methods are based on Wasser, et al. (2004)
- Isotopic assignment methods are based on Wunder, et al. (2005)

Rundel, C.W., et al. (2013) Novel statistical methods for integrating genetic and stable isotopic data to infer individual-level migratory connectivity. *Molecular Ecology* 22 (16).

Data - DNA microsatellites and $\delta^2\text{H}$

Hermit Thrush (*Catharus guttatus*)

- 138 individuals
- 14 locations
- 6 loci
- 9-27 alleles / locus



Colin Rundel (Duke)

Wilson's Warbler (*Wilsonia pusilla*)

- 163 individuals
- 8 locations
- 9 loci
- 15-31 alleles / locus



© Glenn Bartley

Allele Frequency Model

For the allele i , from locus l , at location k

$$\mathbf{y}_{\cdot lk} | \Theta \sim \text{Mult}(\sum_i y_{ilk}, \mathbf{f}_{lk})$$

$$f_{ilk} = \frac{\exp(\Theta_{ilk})}{\sum_i \exp(\Theta_{ilk})}$$

$$\Theta_{il} | \alpha, \mu \sim \mathcal{N}(\mu_{il}, \Sigma)$$

$$\{\Sigma\}_{i,j} = \alpha_0 \exp\left(- (d_{i,j}/\alpha_1)^{\alpha_2}\right) + \alpha_3 I_{d_{i,j}=0}$$

Genetic Assignment Model

Assignment model:

$$P(S_G | \mathbf{f}, k) = \prod_I P(i_I, j_I | \mathbf{f}, k)$$

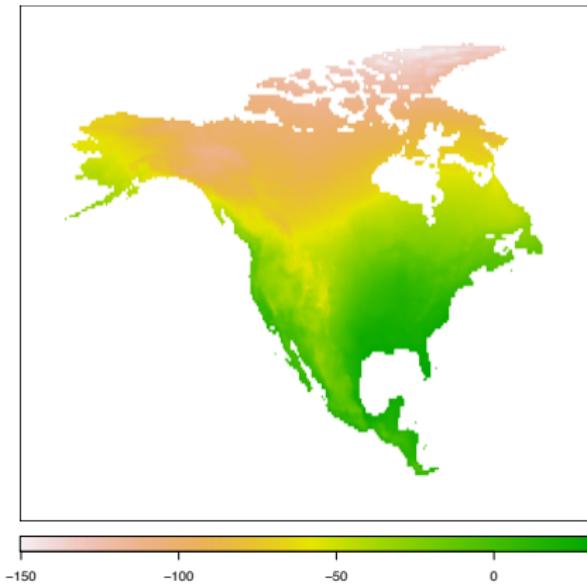
$$P(i_I, j_I | \mathbf{f}, k) = \begin{cases} \gamma P(i_I | \mathbf{f}, k) + (1 - \gamma) P(i_I | \tilde{\mathbf{f}}, k)^2 & \text{if } i = j \\ (1 - \gamma) P(i_I | \mathbf{f}, k) P(j_I | \mathbf{f}, k) & \text{if } i \neq j \end{cases}$$

$$P(i_k | \mathbf{f}, k) = (1 - \delta) f_{lik} + \delta / m_l$$

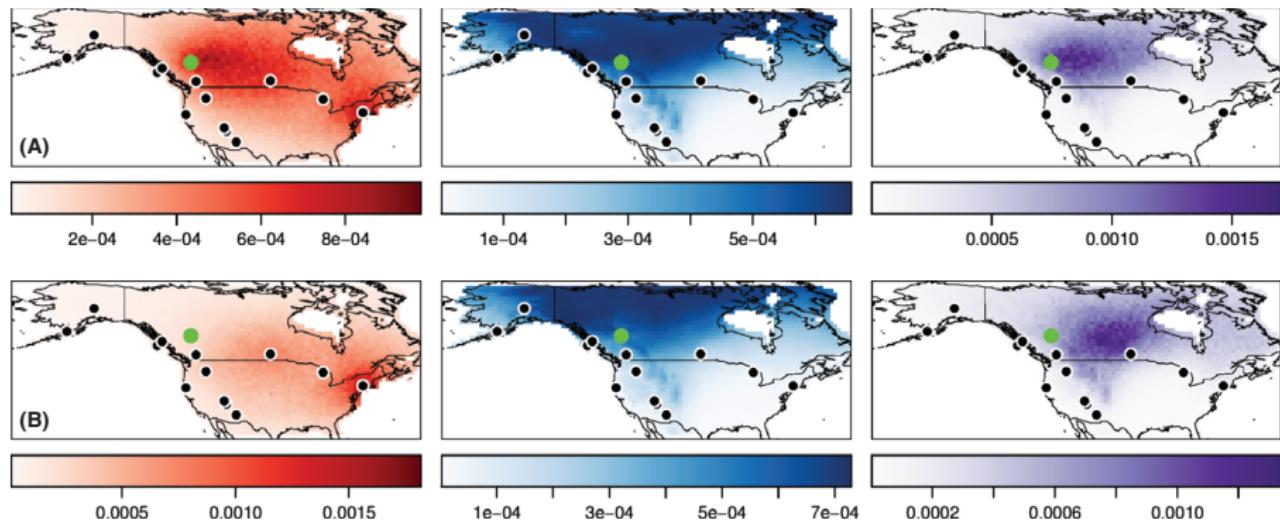
δ – Genotyping error rate γ – Single amplification rate

Isotope Model

$$S_I|k, \tilde{\boldsymbol{p}}, \omega, \rho, \tau^2 \sim N(\omega + \rho \tilde{p}_k, \tau^2)$$



Combined Model

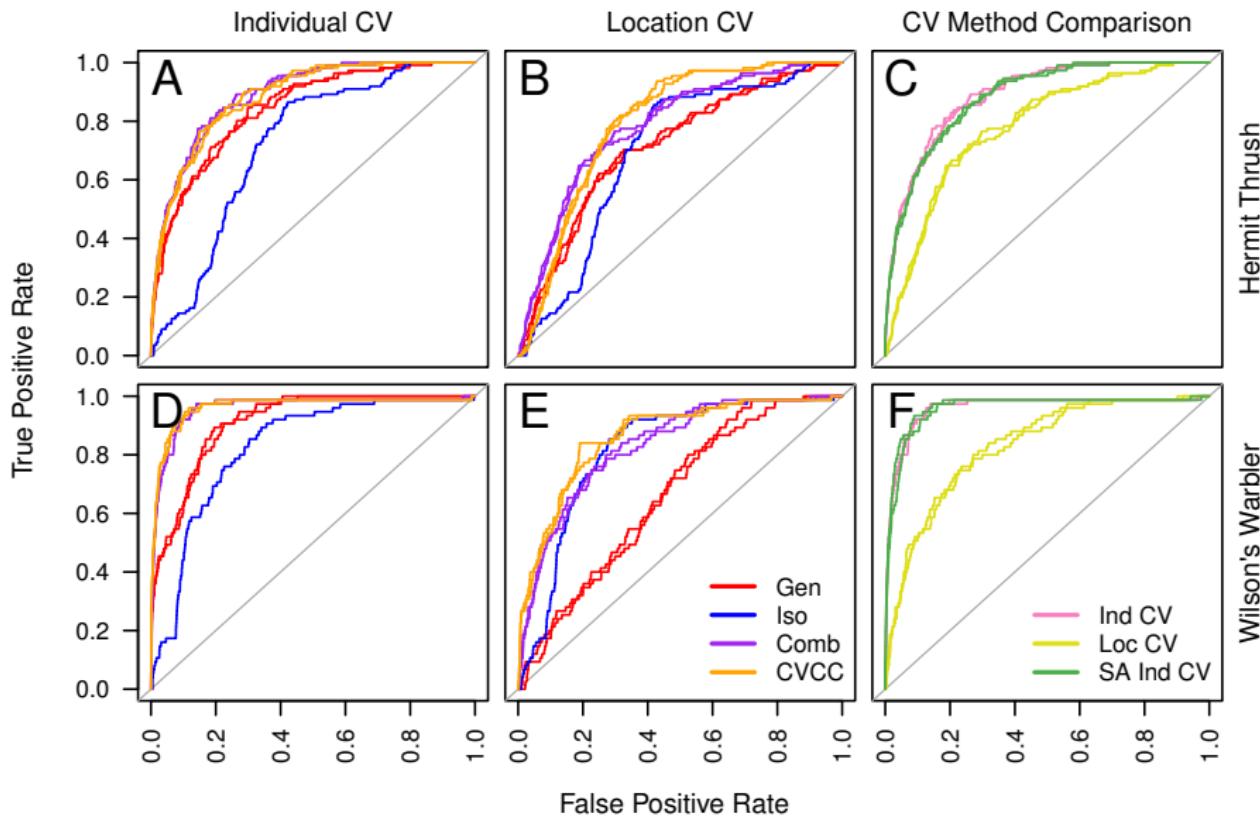


Implementation

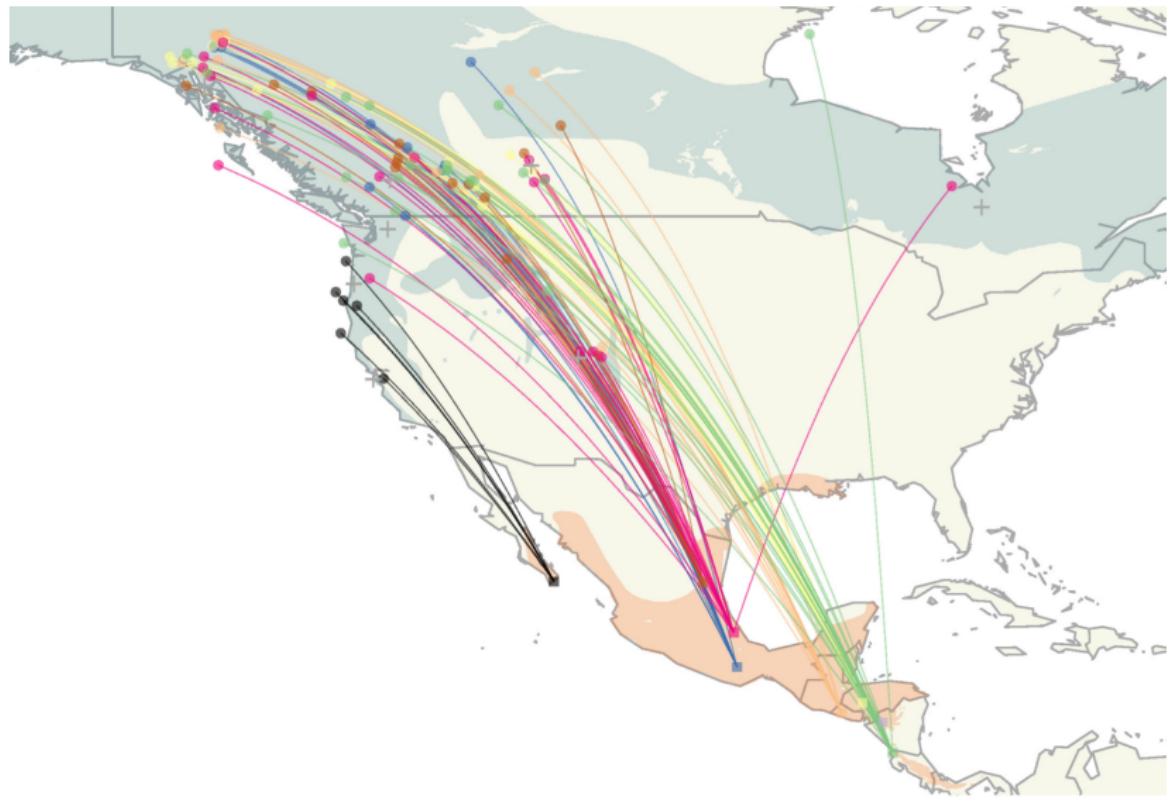
Model fitting is done via MCMC (MH within Gibbs)

- Original implementation in pure C++ with minimal dependencies (Wasser, et al. (2004))
- Rewritten using R / C++ via Rcpp(Armadillo)
 - Code closer to matrix notation (and R)
 - Transparent use of high performance LAPACK implementations
 - R Package - isoscatR - <https://github.com/rundel/isoscatR>
- Model fitting performance is quite good
 - 300,000 iterations in ~ 5.5 minutes
- Bottleneck is sampling posterior predictive
 - 1,000 iterations in ~ 30 minutes

Model Assessment



Migratory Connectivity



Prediction details

Why is the prediction slow?

Prediction details

Why is the prediction slow?

Predicting allele frequencies for Hermit thrush at 3318 novel locations.

To do so we sample from:

$$\Theta_p | \Theta_m \sim \mathcal{N}(\mu_p + \Sigma_{pm} \Sigma_m^{-1} (\Theta_m - \mu_m), \Sigma_p - \Sigma_{pm} \Sigma_m^{-1} \Sigma_{mp})$$

Prediction details

Why is the prediction slow?

Predicting allele frequencies for Hermit thrush at 3318 novel locations.

To do so we sample from:

$$\Theta_p | \Theta_m \sim \mathcal{N}(\mu_p + \Sigma_{pm}\Sigma_m^{-1}(\Theta_m - \mu_m), \Sigma_p - \Sigma_{pm}\Sigma_m^{-1}\Sigma_{mp})$$

Algorithm steps

- ① Calculate Σ_{pm} , Σ_p , and $\Sigma_p - \Sigma_{pm}\Sigma_m^{-1}\Sigma_{mp}$
- ② Calculate $\text{Chol}(\Sigma_p - \Sigma_{pm}\Sigma_m^{-1}\Sigma_{mp})$
- ③ Sample from MVN
- ④ Calculate allele frequencies

Posterior predictive sampling timings

Step	CPU (secs)	CPU+GPU (secs)	Rel. Performance
1. Covariances	1.080	0.046	23.0
2. Cholesky	0.467	0.208	2.3
3. Samples	0.049	0.052	0.9
4. Allele Freq	0.129	0.127	1.0
Total	1.732	0.465	3.7

Posterior predictive sampling timings

Step	CPU (secs)	CPU+GPU (secs)	Rel. Performance
1. Covariances	1.080	0.046	23.0
2. Cholesky	0.467	0.208	2.3
3. Samples	0.049	0.052	0.9
4. Allele Freq	0.129	0.127	1.0
Total	1.732	0.465	3.7

Total run time:

- CPU - 28.9 minutes
- CPU+GPU - 7.8 minutes

Posterior predictive sampling timings

Step	CPU (secs)	CPU+GPU (secs)	Rel. Performance
1. Covariances	1.080	0.046	23.0
2. Cholesky	0.467	0.208	2.3
3. Samples	0.049	0.052	0.9
4. Allele Freq	0.129	0.127	1.0
Total	1.732	0.465	3.7

Total run time:

- CPU - 28.9 minutes
 - CPU+GPU - 7.8 minutes
- × 166 for Hermit Thrush
 × 179 for Wilson's Warbler

Lessons

- Relatively small changes in one function resulted in 3 - 4x improvement
 - Cross validation results in two days instead of a week
 - 1-2 weeks of implementation, 1 week of tweaking / testing
 - Started with Cholesky, other optimizations followed

Lessons

- Relatively small changes in one function resulted in 3 - 4x improvement
 - Cross validation results in two days instead of a week
 - 1-2 weeks of implementation, 1 week of tweaking / testing
 - Started with Cholesky, other optimizations followed
- Issues
 - External library dependency makes package development (much) more complicated
 - Additional code verbosity and complexity

Improving Covariance Calculations

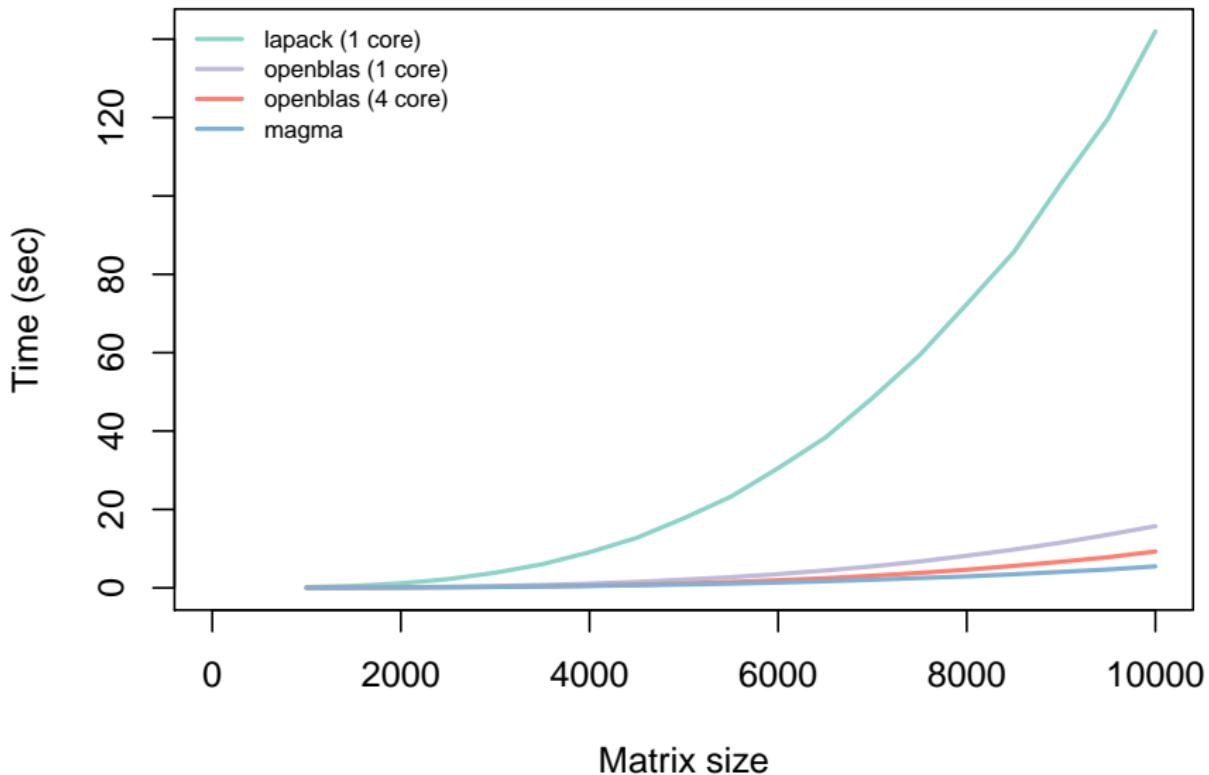
Covariance is assumed to be stationary and isotropic

- Elements of the covariance matrix can be calculated independently
- Small scale “embarrassingly parallel”
- Implementation is straight forward
(if we don't worry about things like symmetry)

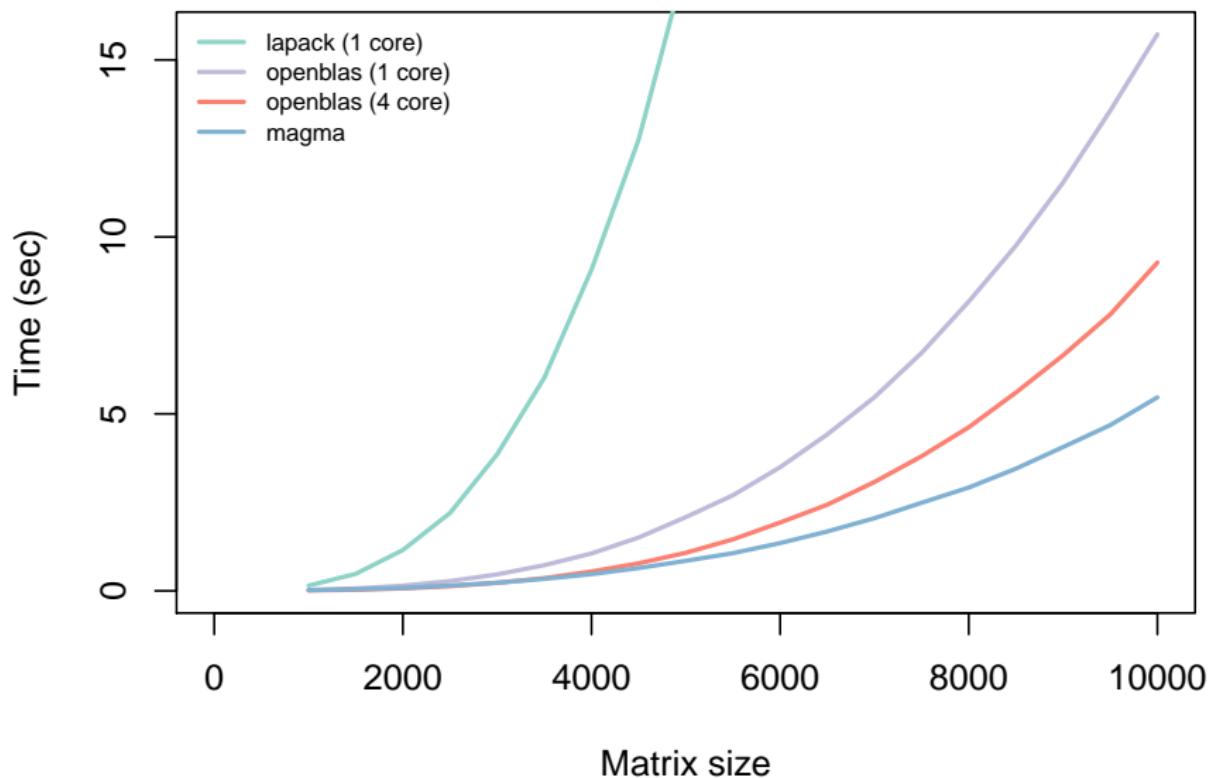
```
__global__ void powered_exponential_cov_kernel(double* dist, double* cov,
                                                const int nm, const double sigma2,
                                                const double phi, const double kappa)
{
    int n_threads = gridDim.x * blockDim.x;
    int pos = blockDim.x * blockIdx.x + threadIdx.x;

    for (int i = pos; i < nm; i += n_threads)
    {
        cov[i] += sigma2 * exp( -pow(dist[i] * phi, kappa) );
    }
}
```

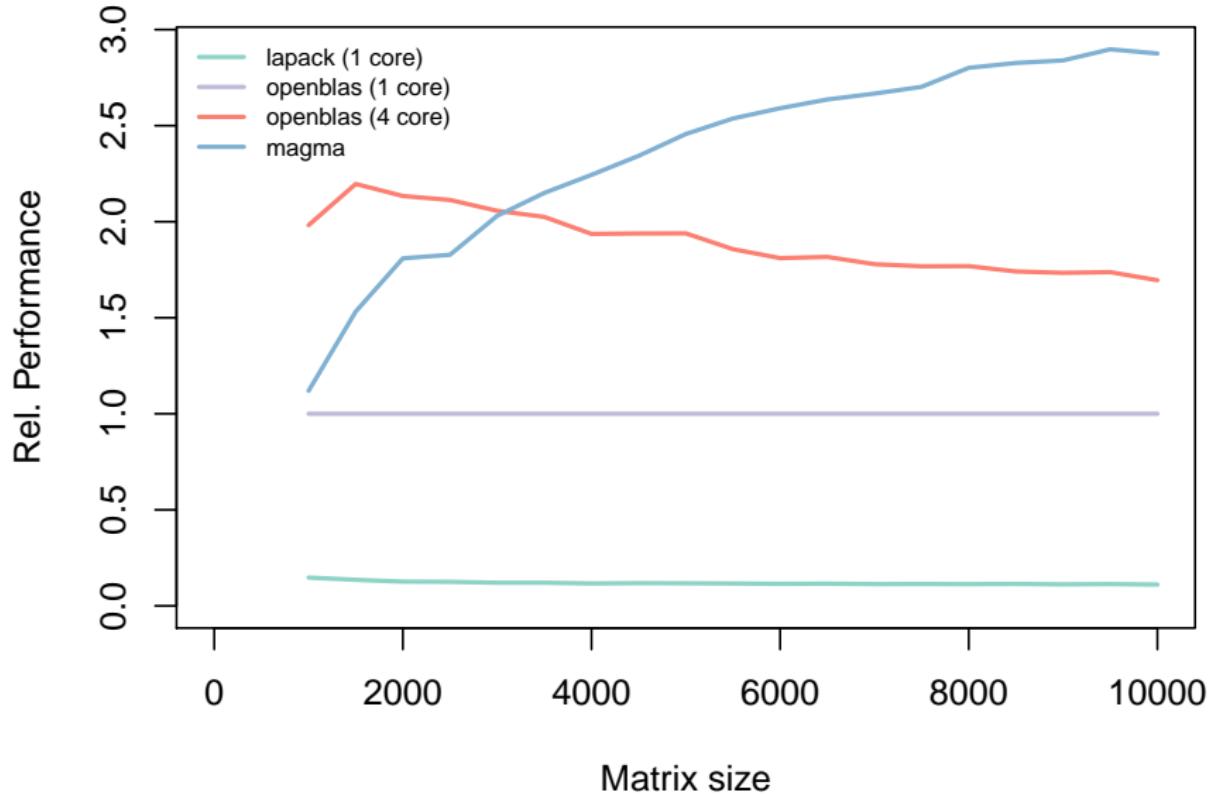
Improving Cholesky



Improving Cholesky



Improving Cholesky



Building core tools

Common set of (expensive) tasks for GP models

- Covariance calculation
- Cholesky of Cov.
- Inverse of Cov.

Goal is to make performing these tasks on a GPU as painless as possible and allow interoperability with GPU (magma, CUBLAS) and CPU (Armadillo) libraries.

- GPU matrix class
- Modern resource management (RAII)
- Simple translation to and from

R Package - RcppGP - <https://github.com/rundel/RcppGP>

CPU vs GPU code

```
arma::mat prop_Sigma = arma::exp(-prop_phi * d_CIF);
arma::mat prop_Sigma_U = arma::chol(prop_Sigma);

double prop_Sigma_log_det = 2*arma::accu(arma::log(prop_Sigma_U.diag()));

arma::mat prop_Sigma_U_inv = arma::inv(arma::trimatu(prop_Sigma_U));
arma::mat prop_Sigma_inv = prop_Sigma_U_inv * prop_Sigma_U_inv.t();
```

```
exponential_cov_gpu(d_CIF_gpu.mat, cov_gpu.mat, nr_CIF, nr_CIF, 1.0, prop_phi, 64);
cov_gpu.chol('L',false);

double prop_Sigma_log_det = 2*arma::accu(arma::log(cov_gpu.get_mat().diag()));

cov_gpu.inv_chol('L',true);
arma::mat prop_Sigma_inv = cov_gpu.get_mat();
```

Back

1 Background

2 Migratory Bird Spatial Assignment Model

3 Speciated PM_{2.5} Modeling

4 Related and Future Work

Background

Fine particulate matter ($\text{PM}_{2.5}$) is an EPA regulated air pollutant linked to a variety of adverse health effects

- Classified based on particle size ($< 2.5 \mu\text{m}$ diameter)
- Major species: Sulfate, Nitrate, Ammonium, Soil, Carbon.
- Minor species: trace elements (K, Mg, Ca), heavy metals (Cu, Fe), etc.
- Complex spatio-temporal dependence between species

Data

Speciated PM_{2.5} Sources

- Chemical Speciation Network (CSN) - 221 stations
- Interagency Monitoring of Protected Visual Environments (IMPROVE) - 172 stations

Total PM_{2.5} Sources

- Federal Reference Method (FRM) - 949 stations

Model Output

- Community Multi-scale Air Quality (CMAQ) - 12 km grid

Data Issues

- Monitoring frequency
- Total vs Sum of Species

Model Details

For each species and network,

$$C_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{C,t}^i(\mathbf{s})$$

$$Z_t^i(\mathbf{s}) = \max(0, \tilde{Z}_t^i(\mathbf{s}))$$

$$I_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{I,t}^i(\mathbf{s})$$

$$\tilde{Z}_t^i(\mathbf{s}) = \beta_{0,t}^i + \beta_{0,t}^i(\mathbf{s}) + \beta_{1,t}^i Q_t^i(B_{\mathbf{s}})$$

Model Details

For each species and network,

$$C_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{C,t}^i(\mathbf{s})$$

$$Z_t^i(\mathbf{s}) = \max(0, \tilde{Z}_t^i(\mathbf{s}))$$

$$I_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{I,t}^i(\mathbf{s})$$

$$\tilde{Z}_t^i(\mathbf{s}) = \beta_{0,t}^i + \beta_{0,t}^i(\mathbf{s}) + \beta_{1,t}^i Q_t^i(B_{\mathbf{s}})$$

For total PM,

$$C_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{C,t}^{tot}(\mathbf{s})$$

$$Z_t^{tot}(\mathbf{s}) = \sum_{i=1}^5 Z_t^i(\mathbf{s}) + Z_t^o(\mathbf{s})$$

$$I_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{I,t}^{tot}(\mathbf{s})$$

$$F_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{F,t}^{tot}(\mathbf{s})$$

Model Details

For each species and network,

$$C_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{C,t}^i(\mathbf{s})$$

$$Z_t^i(\mathbf{s}) = \max(0, \tilde{Z}_t^i(\mathbf{s}))$$

$$I_t^i(\mathbf{s}) = Z_t^i(\mathbf{s}) + \epsilon_{I,t}^i(\mathbf{s})$$

$$\tilde{Z}_t^i(\mathbf{s}) = \beta_{0,t}^i + \beta_{0,t}^i(\mathbf{s}) + \beta_{1,t}^i Q_t^i(B_{\mathbf{s}})$$

For total PM,

$$C_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{C,t}^{tot}(\mathbf{s})$$

$$Z_t^{tot}(\mathbf{s}) = \sum_{i=1}^5 Z_t^i(\mathbf{s}) + Z_t^o(\mathbf{s})$$

$$I_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{I,t}^{tot}(\mathbf{s})$$

$$F_t^{tot}(\mathbf{s}) = Z_t^{tot}(\mathbf{s}) + \epsilon_{F,t}^{tot}(\mathbf{s})$$

$$Z_t^o(\mathbf{s}) = \max(0, \tilde{Z}_t^o(\mathbf{s}))$$

$$\tilde{Z}_t^o(\mathbf{s}) = \beta_{0,t}^o + \beta_{0,t}^o(\mathbf{s}) + \beta_{1,t}^o Q_t^o(B_{\mathbf{s}})$$

Model Details Cont.

Spatial dependence enters the model through the $\beta_{0,t}^i(s)$ and $\beta_{0,t}^o(s)$ parameters.

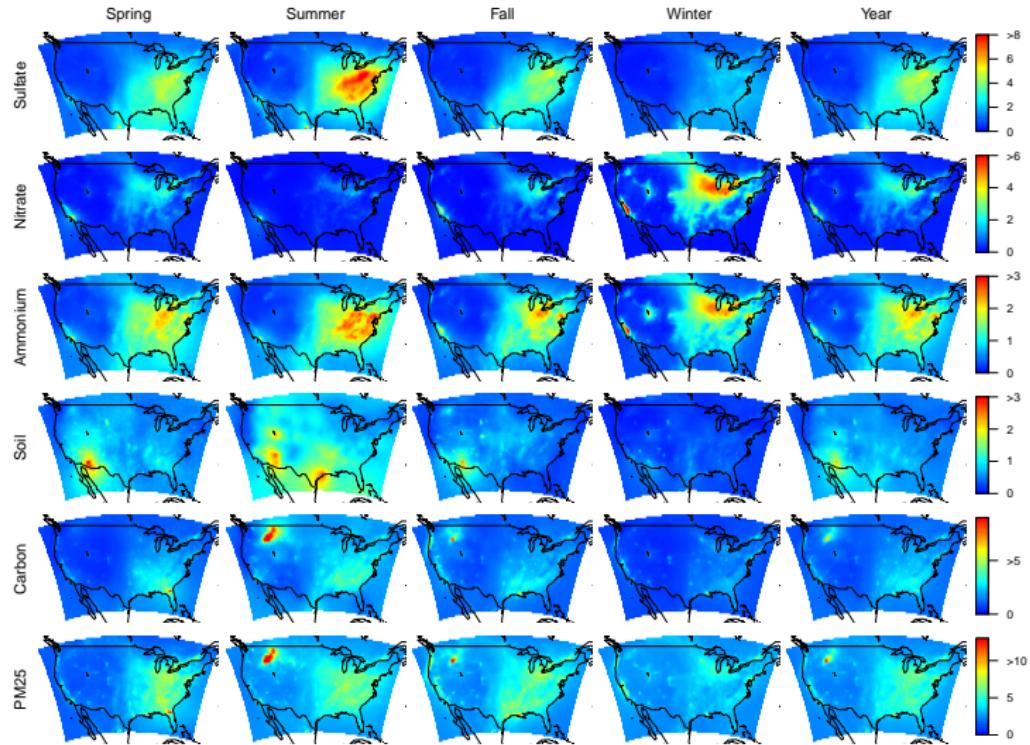
$$\beta_{0,t}^i(\mathbf{s}) = \sigma_t^i w_t^i(\mathbf{s}) \quad \beta_{0,t}^o(\mathbf{s}) = \eta_t w_t^o(\mathbf{s})$$

where $w_t^i(\mathbf{s})$ and $w_t^o(\mathbf{s})$ are zero mean, variance 1, Gaussian processes with exponential correlation given by

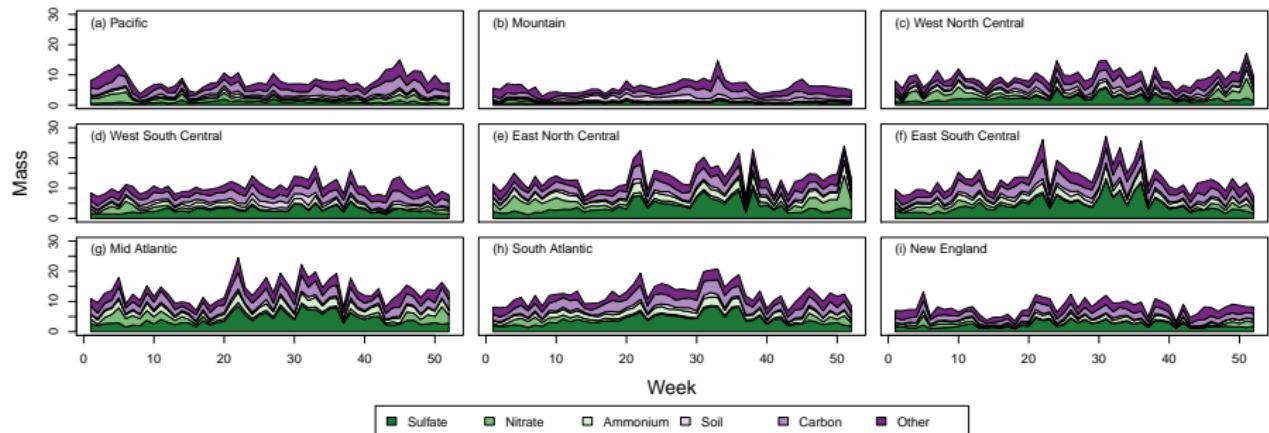
$$\text{corr}(w_t^i(\mathbf{s}), w_t^i(\mathbf{s}')) = \exp(-\phi_t^i |\mathbf{s} - \mathbf{s}'|)$$

$$\text{corr}(w_t^o(\mathbf{s}), w_t^o(\mathbf{s}')) = \exp(-\xi_t |\mathbf{s} - \mathbf{s}'|)$$

Model results



Model results



Model fitting performance

Parameter	CPU (secs)	CPU+GPU (secs)	Rel. Performance
β_0, β_1	0.00029	0.00030	0.97
$\beta_0(s)$	0.09205	0.09132	1.00
σ^2, η^2	0.00383	0.00385	0.99
ϕ, ξ	0.46084	0.25174	1.83
τ_i^2, τ_{tot}^2	0.00003	0.00003	1.00
Total	0.55708	0.34729	1.60

Run times

Total run time for model fitting (50,000 iterations):

- CPU - 7.7 hours \times 52 weeks
- CPU+GPU - 4.8 hours

Run times

Total run time for model fitting (50,000 iterations):

- CPU - 7.7 hours × 52 weeks
- CPU+GPU - 4.8 hours

Total run time for model prediction at 5950 locations (1,000 iterations):

- CPU - 7.2 hours × 52 weeks
- CPU+GPU - 4.3 hours

Run times

Total run time for model fitting (50,000 iterations):

- CPU - 7.7 hours × 52 weeks
- CPU+GPU - 4.8 hours

Total run time for model prediction at 5950 locations (1,000 iterations):

- CPU - 7.2 hours × 52 weeks
- CPU+GPU - 4.3 hours

One run takes about 775 hours total on CPU alone, 473 on CPU and GPU.

Lessons

- Infrastructure makes a huge difference in development time
 - 1 hour to go from CPU implementation to CPU+GPU implementation
 - Code shown previously is 2/3 of the changes necessary [Code](#)
- In practice, easier to run CPU only code across more servers (configuration time / effort)
 - Not possible (or at least easy) for models variants that are not independent in time.
- Rcpp Attributes offer huge advantages in development and deployment
 - Simplifies external dependencies (locally)

1 Background

2 Migratory Bird Spatial Assignment Model

3 Speciated PM_{2.5} Modeling

4 Related and Future Work

GPUs and reduced rank GPs

For example the Gaussian Process,

$$Y(\mathbf{s}) = \mathbf{x}'(\mathbf{s})\boldsymbol{\beta} + w(\mathbf{s}) + \epsilon(\mathbf{s}), \quad \epsilon(\mathbf{s}) \sim N(0, \tau^2 I_n)$$
$$w(\mathbf{s}) \sim N(0, \mathbf{C}(\mathbf{s})), \quad \mathbf{C}(\mathbf{s}, \mathbf{s}') = \sigma^2 \rho(\mathbf{s}, \mathbf{s}' | \theta)$$

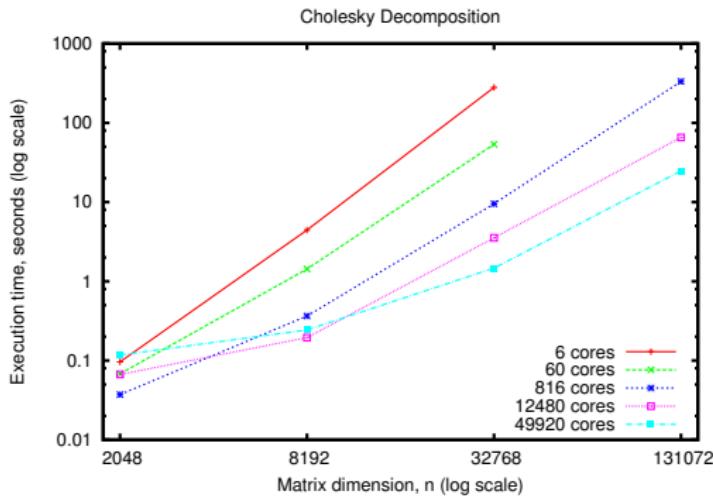
can be approximated using a Gaussian predictive process with knots at locations \mathbf{s}^* where

$$Y(\mathbf{s}) = \mathbf{x}'(\mathbf{s})\boldsymbol{\beta} + \tilde{w}(\mathbf{s}) + \epsilon(\mathbf{s})$$
$$\tilde{w}(\mathbf{s}) = \mathbf{C}(\mathbf{s}, \mathbf{s}^*) \mathbf{C}(\mathbf{s}^*)^{-1} w(\mathbf{s}^*).$$
$$w(\mathbf{s}^*) \sim N(0, \mathbf{C}(\mathbf{s}^*))$$

Another Approach

bigGP is an R package written by Chris Paciorek, et al.

- Specialized implementation of LA operation for GPs
- Designed to run on large super computer clusters
- Uses both shared and distributed memory
- Able to fit models on the order of $n = 65k$ (32 GB Cov. matrix)



Future Directions

- Refinement of RcppGP
 - Full Rcpp Attribute support
 - Transparent GPU to CPU failover
- Single vs. Multi-GPU algorithms
- Mixed precision methods
- NVBLAS
- Unified memory

Acknowledgments

Migratory Connectivity

- John Novembre - UChicago
- Thomas Smith - UCLA
- Kristen Ruegg - UCLA, UCSC
- Center for Tropical Research,
UCLA IoES

Speciated PM_{2.5}

- Alan Gelfand - Duke
- Dave Holland - EPA
- Erin Schliep - Duke

Performance

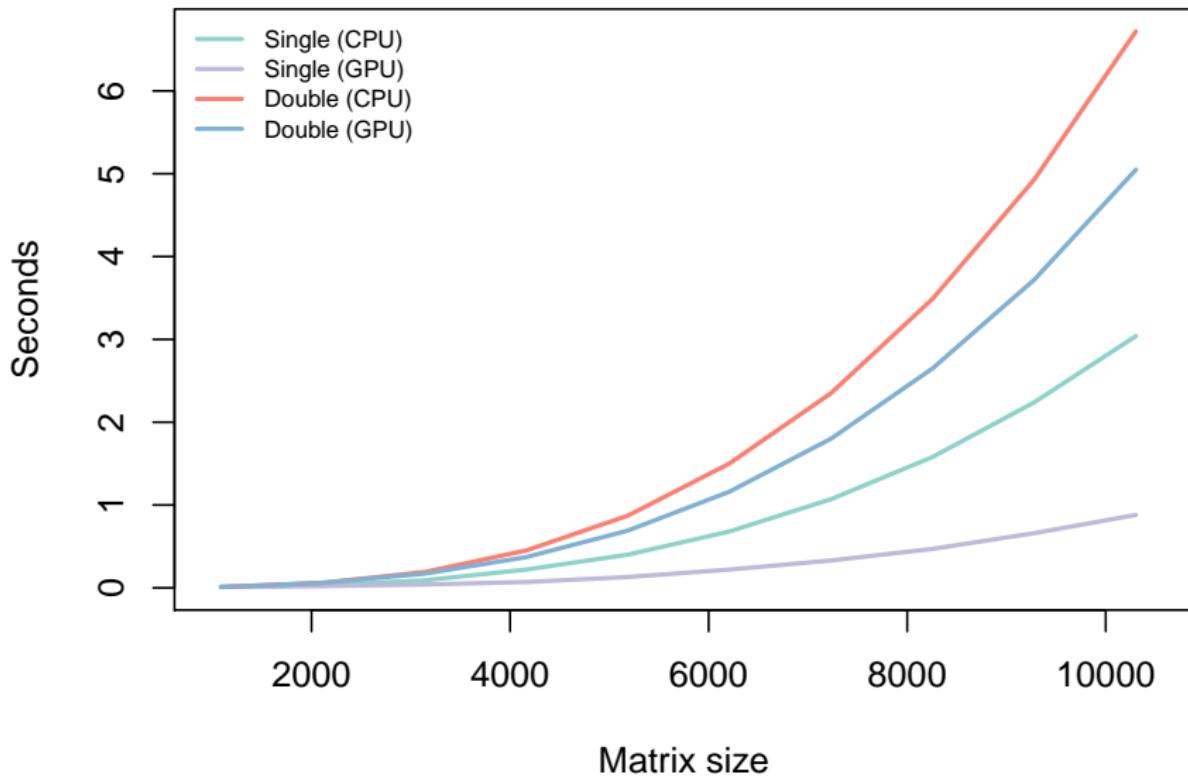
System Specs:

- 4 core Intel i5-2500K @ 3.30 GHz
- 16 GB DDR3 @ 1333 MHz
- GeForce GTX 460

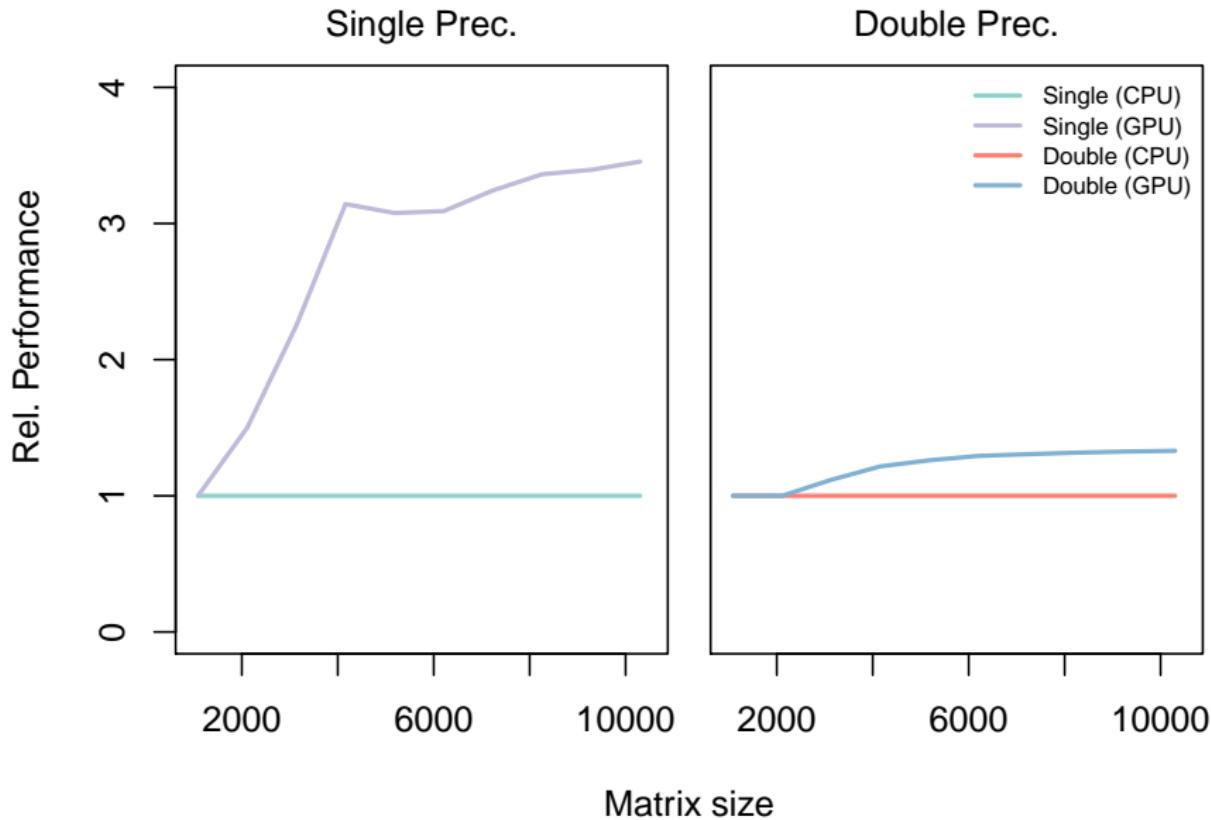
Software Specs:

- Ubuntu 13.10
- OpenBlas 0.2.8
- CUDA 6.0RC
- Magma 1.4.1

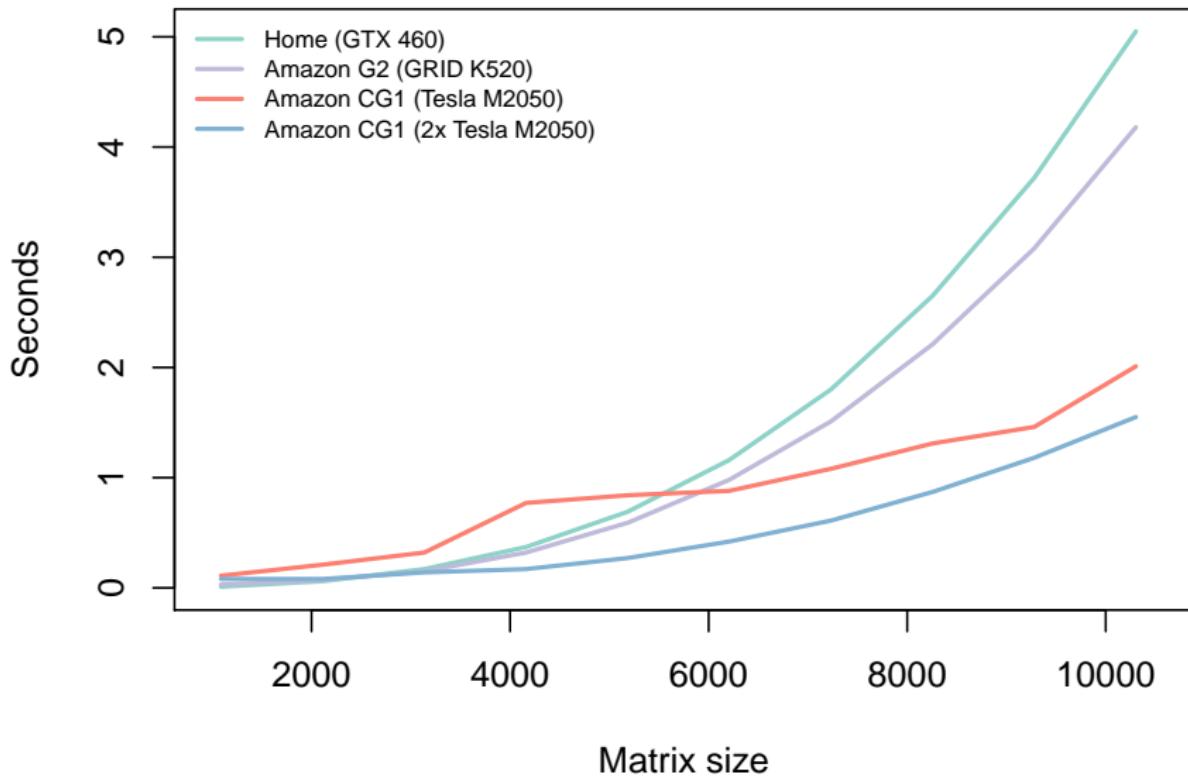
Double vs single float performance



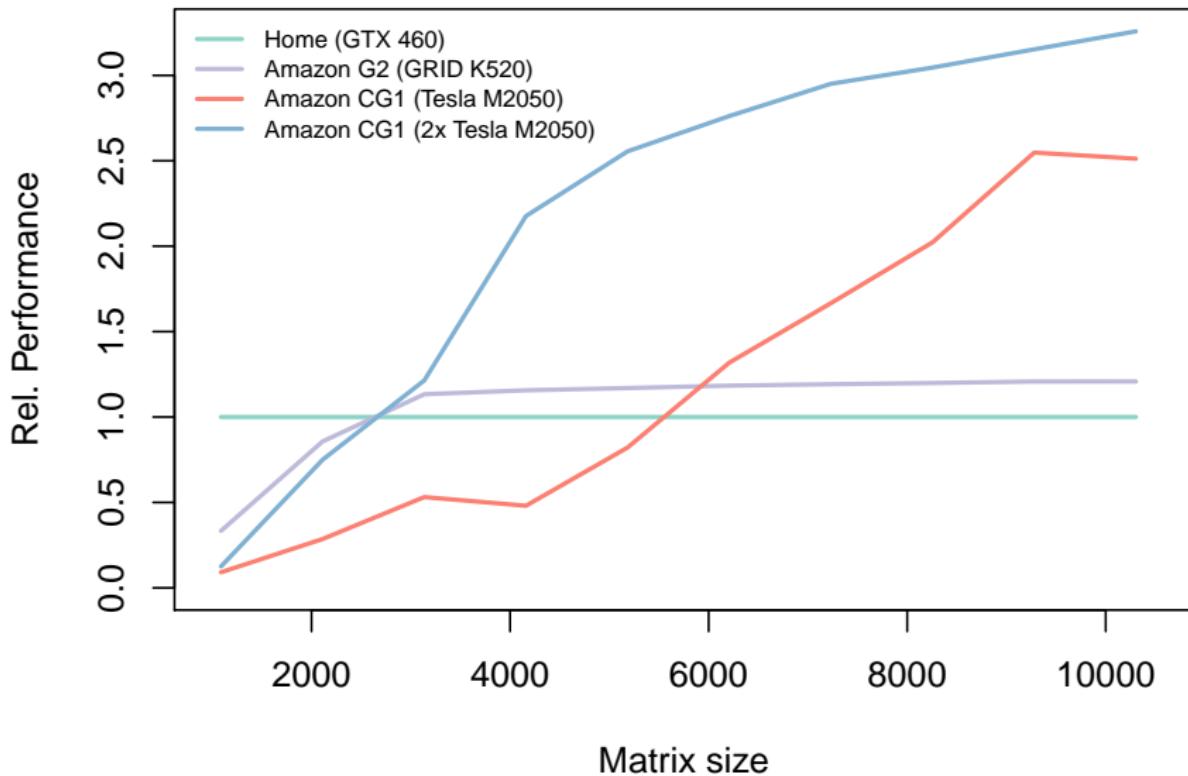
Double vs single float performance



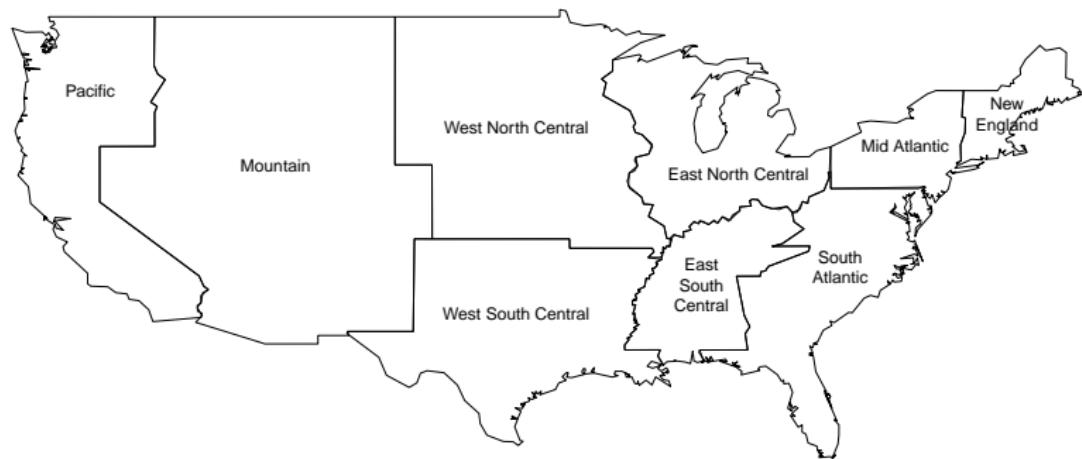
Amazon EC2 Instances



Amazon EC2 Instances



Regions



Back