

An overview of (*some*)  
Bayesian computational frameworks  
for teaching

# Personal Context

What I teach:

- Sta 323 - Statistical Computing (R)
- Sta 523 - Programming for Statistical Science (R)
- STA 663 - Statistical Computing and Computation (Python)
- STA 344/444/644 - Spatio-temporal modeling (R)

Tools I use:

- The majority of my time is spent with R, with a bit of C++
- I use JAGS and Stan for applied modeling
- Recently, more teaching focused on the Python ecosystem

# Bayesian computational frameworks?

A collection of tools that implement a domain specific language for expressing and implementing Bayesian statistical models.

For example,

- JAGS
- STAN
- pymc
- + many others

# Some teaching considerations

- Ease of use (installation, syntax, debugging, etc.)
- Blackboxiness / High vs low level
- Generalizability
- Performance / Limitations
- Wider curriculum

# Installation + basic usage

- All of the frameworks have external / system dependencies
  - e.g. libjags, Eigen, theano, etc.
- Generally easy to install binary packages are available
  - source installs can be challenging
- If things break it tends to be spectacular and difficult to troubleshoot
  - OS makes a difference
  - 🔥 Burn it down 🔥 as a path forward

# Example - Bayesian SLR

```
d = read.csv("data/lm.csv")  
plot(d)
```

$$y_i | m, b, \sigma \sim N(m \cdot x_i + b, \sigma)$$

$$m \sim N(0, 10)$$

$$n \sim N(0, 10)$$

$$\sigma \sim N(0, 5)$$

# SLR - JAGS

```
model = "  
model{  
  m ~ dnorm(0, 1/100)  
  b ~ dnorm(0, 1/100)  
  
  sigma ~ dnorm(0, 1/25) T(0,)  
  
  for(i in 1:length(y)) {  
    mu[i] = m*x[i] + b  
    y[i] ~ dnorm(mu[i], 1/(sigma^2))  
  }  
}  
"
```

```
jags_model = rjags::jags.model(  
  textConnection(model), data = d, n.chains=4  
)
```

```
## Compiling model graph  
##   Resolving undeclared variables  
##   Allocating nodes  
## Graph information:  
##   Observed stochastic nodes: 11  
##   Unobserved stochastic nodes: 3  
##   Total graph size: 56  
##  
## Initializing model
```

```
update(jags_model, n.iter=1000, progress.bar="none")  
  
post_jags = rjags::coda.samples(  
  jags_model, variable.names=c("m","b"),  
  n.iter=1000, progress.bar="none"  
)
```

# SLR - Stan

```
stan = "  
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real m;  
  real b;  
  real<lower=0> sigma;  
}  
transformed parameters {  
  vector[N] mu = m*x + b;  
}  
model {  
  m ~ normal(0, 10);  
  b ~ normal(0, 10);  
  sigma ~ normal(0, 5);  
  y ~ normal(mu, sigma);  
}  
"
```

```
post_stan = rstan::stan(  
  model_code = stan,  
  data = list(x=d$x, y=d$y, N=nrow(d)),  
  pars = c("m", "b", "sigma"),  
  chains = 4, warmup = 1000, iter = 2000,  
  refresh = 1000, verbose = FALSE,  
)
```

```
##  
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).  
## Chain 1:  
## Chain 1: Gradient evaluation took 1.6e-05 seconds  
## Chain 1: 1000 transitions using 10 leapfrog steps per transition  
## Chain 1: Adjust your expectations accordingly!  
## Chain 1:  
## Chain 1:  
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)  
## Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)  
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)  
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)  
## Chain 1:  
## Chain 1: Elapsed Time: 0.012 seconds (Warm-up)  
## Chain 1:                0.012 seconds (Sampling)  
## Chain 1:                0.024 seconds (Total)  
## Chain 1:  
##  
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).  
## Chain 2:  
## Chain 2: Gradient evaluation took 1e-06 seconds  
## Chain 2: 1000 transitions using 10 leapfrog steps per transition  
## Chain 2: Adjust your expectations accordingly!
```



# SLR - pymc3

```
import pymc3 as pm
import arviz as az
```

```
with pm.Model() as lm:
    m = pm.Normal('m', mu=0, sd=10)
    b = pm.Normal('b', mu=0, sd=10)

    mu = m * d.x + b
    sigma = pm.HalfNormal('sigma', sd=5)

    y = pm.Normal('y', mu=mu, sd=sigma, observed=d.y)
```

```
with lm:
    post_pymc = pm.sample(return_inferencedata=True, random_seed=
```

```
## ■
## Auto-assigning NUTS sampler...
## Initializing NUTS using jitter+adapt_diag...
## Multiprocess sampling (4 chains in 4 jobs)
## NUTS: [sigma, b, m]
## Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000
## There were 6 divergences after tuning. Increase `target_accept`
```

# Modelling results

- JAGS models return a coda `mcmc.list`
  - basic tabular structure that is easy to work with
- Stan models return a `stanfit` S4 object
  - less directly accessible but provides important basic summaries (e.g. `n_eff`, `Rhat`, etc.)
  - easily convertible to coda (`As.mcmc.list()`)
- pymc3 models return\* ArviZ `InferenceData` objects (xarray/NetCDF based)
  - complex schema (everything and the kitchen sink approach)
  - less tabular friendly
- All frameworks support quick basic visualizations of results (trace, density, caterpillar, etc.)

# Error reporting

- As pymc models are Python code, any syntax errors are reported as Python syntax errors
- JAGS and Stan implement their own parsers which have generally helpful error messages with the former tending to be terser / less detailed,

```
rjags::jags.model(  
  textConnection("  
  model{  
    m ~ dnorm(0, 1/100  
  }  
  ")  
)
```

```
## Error in rjags::jags.model(textConnection("\n  model{\n  m  
## Error parsing model file:  
## syntax error on line 4 near "}"
```

```
rstan::stan(  
  model_code = "  
  model {  
    m ~ normal(0, 10;  
  }  
  "  
)
```

```
## Error in stanc(file = file, model_code = model_code, model_n  
## 0  
##  
## Syntax error in 'string', line 2, column 20 to column 21,  
## parsing error:  
##  
## Ill-formed phrase. Found an expression. This can be followed  
## by a ",", a "}", a ")", a "[", a "]" or an infix or postfi  
## operator.
```

- Runtime errors are a mixed bag

# Posterior predictive checks

- Possible with all three frameworks, JAGS and Stan require that extra parameters be included in the model:

JAGS:

```
for(i in 1:length(y)) {  
  y_tilde[i] ~ dnorm(mu[i], 1/(sigma^2))  
}
```

Stan:

```
generated quantities {  
  real y_tilde[N] = normal_rng(mu, sigma);  
}
```

- pymc allows for the PPD to be sampled from an existing model result,

```
with lm:  
  y_tilde = pm.sample_posterior_predictive(  
    post_pymc, var_names=["y"], random_seed=1234  
  )
```

- Similarly, the prior predictive samples can be generated without rewriting the model

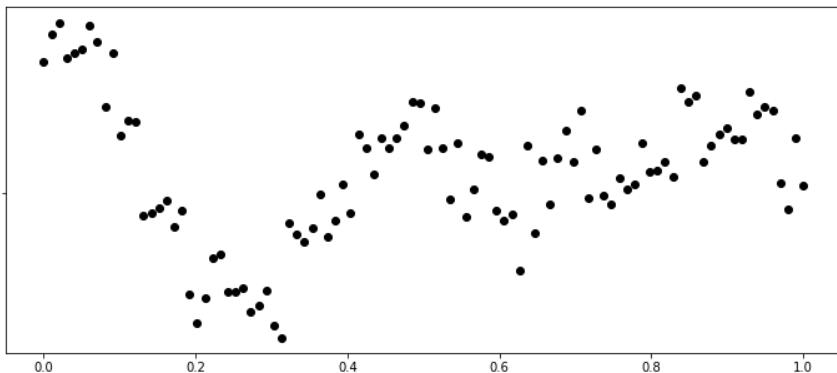
# Limitations - GP Reg

```
d = pd.read_csv("data/gp.csv")
d.shape
```

```
## (100, 3)
```

```
D = np.array([ np.abs(xi - d.x) for xi in d.x])
I = (D == 0).astype("double")
```

```
fig = plt.figure(figsize=(12, 5))
plt.plot(d.x, d.y, "ok", ".")
plt.show()
```



```
with pm.Model() as gp:
    nugget = pm.HalfCauchy("nugget", beta=5)
    sigma2 = pm.HalfCauchy("sigma2", beta=5)
    ls      = pm.HalfCauchy("ls",      beta=5)

    Sigma = I * nugget + sigma2 * np.exp(-0.5 * D**2 * ls**2)

    y = pm.MvNormal(
        "y",
        mu=np.zeros(d.shape[0]),
        cov=Sigma, observed=d.y
    )
```

# NUTS

```
with gp:
    post_nuts = pm.sample(
        return_inferencedata = True,
        chains = 2
    )
```

```
## ■
## Multiprocess sampling (2 chains in 4 jobs)
## NUTS: [ls, sigma2, nugget]
## Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 240 seconds.
```

```
az.summary(post)
```

```
##          mean      sd  hdi_3%  hdi_97%  ...  mcse_sd  ess_bulk  ess_tail  r_hat
## nugget    0.541  0.087   0.397   0.715  ...    0.002   1754.0   1292.0    1.0
## sigma2    4.096  2.557   1.262   8.273  ...    0.060   1067.0   1004.0    1.0
## ls       10.756  2.383   6.593  15.267  ...    0.049   1068.0   1109.0    1.0
##
## [3 rows x 9 columns]
```

# Slice steps

```
with gp:
    step = pm.Slice([nugget, sigma2, ls])
    post_slice = pm.sample(
        return_inferencedata = True,
        chains = 2,
        step = step
    )
```

```
## ■
## Multiprocess sampling (2 chains in 4 jobs)
## CompoundStep
## >Slice: [ls]
## >Slice: [sigma2]
## >Slice: [nugget]
## Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 24 seconds.
```

```
az.summary(post_slice)
```

```
##          mean      sd  hdi_3%  hdi_97%  ...  mcse_sd  ess_bulk  ess_tail  r_hat
## nugget    0.542  0.085   0.399   0.705  ...    0.002   1573.0   1510.0    1.0
## sigma2    4.557  3.551   1.082  10.070  ...    0.087    915.0    842.0    1.0
## ls       10.526  2.466   5.815  14.552  ...    0.055    989.0    967.0    1.0
##
## [3 rows x 9 columns]
```

# Metropolis-Hastings steps

```
with gp:
    step = pm.Metropolis([nugget, sigma2, ls])
    post_mh = pm.sample(
        return_inferencedata = True,
        chains = 2,
        step = step
    )
```

```
## ■
## Multiprocess sampling (2 chains in 4 jobs)
## CompoundStep
## >Metropolis: [ls]
## >Metropolis: [sigma2]
## >Metropolis: [nugget]
## Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 9 seconds.
## The estimated number of effective samples is smaller than 200 for some parameters.
```

```
az.summary(post_mh)
```

```
##          mean      sd  hdi_3%  hdi_97%  ...  mcse_sd  ess_bulk  ess_tail  r_hat
## nugget    0.546  0.096   0.373    0.722  ...    0.004    321.0    351.0    1.01
## sigma2    4.535  3.522   1.081    9.282  ...    0.155    231.0    273.0    1.03
## ls       10.518  2.314   6.730   14.987  ...    0.118    188.0    220.0    1.03
##
## [3 rows x 9 columns]
```



# Concluding thoughts

- All of these frameworks are a reasonable choice
  - Many different axes to optimize over
  - More excellent choices than ever before
- Feeling the grass is greener is real
- Personal choice for Fall 2022 (Sta 344)?
  - Probably BRMS -> Stan

# Thank you!

## Questions or Comments?



[rundel.github.io](https://rundel.github.io)



[rundel](https://github.com/rundel)



[rundel](https://twitter.com/rundel)



[rundel@gmail.com](mailto:rundel@gmail.com)  
[colin.rundel@duke.edu](mailto:colin.rundel@duke.edu)



[bit.ly/rundel\\_isba2022](https://bit.ly/rundel_isba2022)

---